



## 50 Startups Predictions using Regression algorithms

```
In [1]: import warnings
warnings.filterwarnings('ignore')
import pandas as pd
import numpy as np
df = pd.read_csv("E:/Exposys Data labs/50_Startups.csv")
df
```

Out[1]:

	R&D Spend	Administration	Marketing Spend	Profit
0	165349.20	136897.80	471784.10	192261.83
1	162597.70	151377.59	443898.53	191792.06
2	153441.51	101145.55	407934.54	191050.39
3	144372.41	118671.85	383199.62	182901.99
4	142107.34	91391.77	366168.42	166187.94
5	131876.90	99814.71	362861.36	156991.12
6	134615.46	147198.87	127716.82	156122.51
7	130298.13	145530.06	323876.68	155752.60
8	120542.52	148718.95	311613.29	152211.77
9	123334.88	108679.17	304981.62	149759.96
10	101913.08	110594.11	229160.95	146121.95
11	100671.96	91790.61	249744.55	144259.40
12	93863.75	127320.38	249839.44	141585.52
13	91992.39	135495.07	252664.93	134307.35
14	119943.24	156547.42	256512.92	132602.65
15	114523.61	122616.84	261776.23	129917.04
16	78013.11	121597.55	264346.06	126992.93
17	94657.16	145077.58	282574.31	125370.37
18	91749.16	114175.79	294919.57	124266.90
19	86419.70	153514.11	0.00	122776.86
20	76253.86	113867.30	298664.47	118474.03
21	78389.47	153773.43	299737.29	111313.02
22	73994.56	122782.75	303319.26	110352.25
23	67532.53	105751.03	304768.73	108733.99
24	77044.01	99281.34	140574.81	108552.04
25	64664.71	139553.16	137962.62	107404.34

	<b>R&amp;D Spend</b>	<b>Administration</b>	<b>Marketing Spend</b>	<b>Profit</b>
<b>26</b>	75328.87	144135.98	134050.07	105733.54
<b>27</b>	72107.60	127864.55	353183.81	105008.31
<b>28</b>	66051.52	182645.56	118148.20	103282.38
<b>29</b>	65605.48	153032.06	107138.38	101004.64
<b>30</b>	61994.48	115641.28	91131.24	99937.59
<b>31</b>	61136.38	152701.92	88218.23	97483.56
<b>32</b>	63408.86	129219.61	46085.25	97427.84
<b>33</b>	55493.95	103057.49	214634.81	96778.92
<b>34</b>	46426.07	157693.92	210797.67	96712.80
<b>35</b>	46014.02	85047.44	205517.64	96479.51
<b>36</b>	28663.76	127056.21	201126.82	90708.19
<b>37</b>	44069.95	51283.14	197029.42	89949.14
<b>38</b>	20229.59	65947.93	185265.10	81229.06
<b>39</b>	38558.51	82982.09	174999.30	81005.76
<b>40</b>	28754.33	118546.05	172795.67	78239.91
<b>41</b>	27892.92	84710.77	164470.71	77798.83
<b>42</b>	23640.93	96189.63	148001.11	71498.49
<b>43</b>	15505.73	127382.30	35534.17	69758.98
<b>44</b>	22177.74	154806.14	28334.72	65200.33
<b>45</b>	1000.23	124153.04	1903.93	64926.08
<b>46</b>	1315.46	115816.21	297114.46	49490.75
<b>47</b>	0.00	135426.92	0.00	42559.73
<b>48</b>	542.05	51743.15	0.00	35673.41
<b>49</b>	0.00	116983.80	45173.06	14681.40

In [2]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   R&D Spend              50 non-null    float64
1   Administration         50 non-null    float64
2   Marketing Spend        50 non-null    float64
3   Profit                 50 non-null    float64
dtypes: float64(4)
memory usage: 1.7 KB
```

In [3]: df.columns

Out[3]: Index(['R&D Spend', 'Administration', 'Marketing Spend', 'Profit'], dtype='object')

**There are no null values so we can start with splitting the data into train and test dataset.**

In [4]: `X = df[['R&D Spend', 'Administration', 'Marketing Spend']]`  
`y = df[['Profit']]`

In [5]: `print(X.shape)`  
`print(y.shape)`

```
(50, 3)
(50, 1)
```

In [6]: `from sklearn.model_selection import train_test_split`  
`X_train,X_test, y_train ,y_test = train_test_split(X,y)`

```
In [7]: print(X_train.shape)
        print(X_test.shape)
        print(y_train.shape)
        print(y_test.shape)
```

```
(37, 3)
(13, 3)
(37, 1)
(13, 1)
```

```
In [8]: y_test[0:10]
```

Out[8]:

	Profit
9	149759.96
4	166187.94
16	126992.93
34	96712.80
12	141585.52
40	78239.91
46	49490.75
8	152211.77
10	146121.95
38	81229.06

## Linear Regression Algorithm

```
In [9]: from sklearn.linear_model import LinearRegression  
lin = LinearRegression()  
lin.fit(X_train,y_train)
```

```
Out[9]: ▾ LinearRegression  
LinearRegression()
```

```
In [10]: print("Coefficient:",lin.coef_)  
print("Intercept:",lin.intercept_)
```

```
Coefficient: [[ 0.84158178 -0.01151763  0.01916807]]  
Intercept: [46985.06267155]
```

```
In [11]: predicted_values = lin.predict(X_test)  
predicted_values[0:10]
```

```
Out[11]: array([[155375.63258322],  
                [172546.13563278],  
                [116305.96240967],  
                [ 88280.72112345],  
                [129301.59483285],  
                [ 73130.97253368],  
                [ 52453.31213318],  
                [152691.58614135],  
                [135872.04456265],  
                [ 66801.52741019]])
```

```
In [12]: print("Train score:",lin.score(X_train,y_train))  
print("Test Score:",lin.score(X_test,y_test))
```

```
Train score: 0.9539670032442588  
Test Score: 0.9263873319485134
```

```
In [13]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
MSE_LR = mean_squared_error(y_test, predicted_values)
MAE_LR = mean_squared_error(y_test, predicted_values)
RMSE_LR = np.sqrt(MSE_LR)
```

```
In [14]: print("Mean squared error for Linear regression algorithm:", MSE_LR)
print("Mean Absolute error for Linear regression algorithm:", MAE_LR)
print("Root Mean Squared error for Linear regression algorithm:", RMSE_LR)
```

Mean squared error for Linear regression algorithm: 88568740.6283923  
Mean Absolute error for Linear regression algorithm: 88568740.6283923  
Root Mean Squared error for Linear regression algorithm: 9411.096675116683

```
In [15]: LR_score = r2_score(y_test, predicted_values)
LR_score
```

Out[15]: 0.9263873319485134

## Ridge Regression Algorithm

```
In [16]: from sklearn.linear_model import Ridge
ridge = Ridge()
ridge.fit(X_train, y_train)
```

Out[16]: A dashed box containing a small downward-pointing triangle followed by the text 'Ridge' and 'Ridge()' on the next line.

```
In [17]: print("Coefficient:", ridge.coef_)
print("Intercept:", ridge.intercept_)
```

Coefficient: [[ 0.84158178 -0.01151763 0.01916807]]  
Intercept: [46985.06267019]



```
In [18]: predicted_values_RR = ridge.predict(X_test)
         predicted_values_RR[0:10]
```

```
Out[18]: array([[155375.63258203],
                [172546.13563124],
                [116305.9624101 ],
                [ 88280.72112515],
                [129301.59483255],
                [ 73130.97253504],
                [ 52453.31213705],
                [152691.58614101],
                [135872.04456151],
                [ 66801.52741119]])
```

```
In [19]: print("Train score:",ridge.score(X_train,y_train))
         print("Test Score:",ridge.score(X_test,y_test))
```

```
Train score: 0.9539670032442586
Test Score: 0.9263873319566358
```

```
In [20]: MSE_RR = mean_squared_error(y_test,predicted_values_RR)
         MAE_RR = mean_squared_error(y_test,predicted_values_RR)
         RMSE_RR = np.sqrt(MSE_LR)
```

```
In [21]: print("Mean squared error for Ridge regression algorithm:",MSE_RR)
         print("Mean Absolute error for Ridge regression algorithm:",MAE_RR)
         print("Root Mean Squared error for Ridge regression algorithm:",RMSE_RR)
```

```
Mean squared error for Ridge regression algorithm: 88568740.61861975
Mean Absolute error for Ridge regression algorithm: 88568740.61861975
Root Mean Squared error for Ridge regression algorithm: 9411.096675116683
```

```
In [22]: RR_score = r2_score(y_test,predicted_values_RR)
         RR_score
```

```
Out[22]: 0.9263873319566358
```

## Lasso Regression Algorithm

```
In [23]: from sklearn.linear_model import Lasso  
lasso = Lasso()  
lasso.fit(X_train,y_train)
```

```
Out[23]: 

▼ Lasso



Lasso()


```

```
In [24]: print("Coefficient:",lasso.coef_)  
print("Intercept:",lasso.intercept_)
```

```
Coefficient: [ 0.84158178 -0.01151763  0.01916807]  
Intercept: [46985.06258987]
```

```
In [25]: predicted_values_LAR = lasso.predict(X_test)  
predicted_values_LAR[0:10]
```

```
Out[25]: array([155375.63261266, 172546.13565212, 116305.96238992,  88280.72109823,  
                129301.59485145,  73130.97247617,  52453.31196028, 152691.58618505,  
                135872.04459675,  66801.52730158])
```

```
In [26]: print("Train score:",lasso.score(X_train,y_train))  
print("Test Score:",lasso.score(X_test,y_test))
```

```
Train score: 0.9539670032442586  
Test Score: 0.9263873314307618
```

```
In [27]: MSE_LAR = mean_squared_error(y_test,predicted_values_LAR)  
MAE_LAR = mean_squared_error(y_test,predicted_values_LAR)  
RMSE_LAR = np.sqrt(MSE_LAR)
```

```
In [28]: print("Mean squared error for Lasso regression algorithm:",MSE_LAR)
print("Mean Absolute error for Lasso regression algorithm:",MAE_LAR)
print("Root Mean Squared error for Lasso regression algorithm:",RMSE_LAR)
```

```
Mean squared error for Lasso regression algorithm: 88568741.2513369
Mean Absolute error for Lasso regression algorithm: 88568741.2513369
Root Mean Squared error for Lasso regression algorithm: 9411.096708212965
```

```
In [29]: LAR_score = r2_score(y_test,predicted_values_LAR)
LAR_score
```

```
Out[29]: 0.9263873314307618
```

## Elastic Net Regression Algorithm

```
In [30]: from sklearn.linear_model import ElasticNet
elastic_net = ElasticNet()
elastic_net.fit(X_train,y_train)
```

```
Out[30]: 

▼ ElasticNet
  ElasticNet()


```

```
In [31]: print("Coefficient:",elastic_net.coef_)
print("Intercept:",elastic_net.intercept_)
```

```
Coefficient: [ 0.84158178 -0.01151763  0.01916807]
Intercept: [46985.06259216]
```

```
In [32]: predicted_values_ENR = elastic_net.predict(X_test)
predicted_values_ENR[0:10]
```

```
Out[32]: array([155375.63256428, 172546.13559882, 116305.96241192,  88280.72115908,
129301.5948336 ,  73130.97254349,  52453.31215675, 152691.58615337,
135872.04454729,  66801.52738436])
```

```
In [33]: print("Train score:",elastic_net.score(X_train,y_train))  
print("Test Score:",elastic_net.score(X_test,y_test))
```

Train score: 0.9539670032442586  
Test Score: 0.9263873319205858

```
In [34]: MSE_ENR = mean_squared_error(y_test,predicted_values_ENR)  
MAE_ENR = mean_squared_error(y_test,predicted_values_ENR)  
RMSE_ENR = np.sqrt(MSE_ENR)
```

```
In [35]: print("Mean squared error for Lasso regression algorithm:",MSE_ENR)  
print("Mean Absolute error for Lasso regression algorithm:",MAE_ENR)  
print("Root Mean Squared error for Lasso regression algorithm:",RMSE_ENR)
```

Mean squared error for Lasso regression algorithm: 88568740.66199419  
Mean Absolute error for Lasso regression algorithm: 88568740.66199419  
Root Mean Squared error for Lasso regression algorithm: 9411.09667690191

```
In [36]: ENR_score = r2_score(y_test,predicted_values_ENR)  
ENR_score
```

Out[36]: 0.9263873319205858

## Support Vector regressor Algorithm

```
In [37]: from sklearn.svm import SVR  
svr = SVR(kernel='linear')  
svr.fit(X_train,y_train)
```

Out[37]:

▼

SVR

SVR(kernel='linear')

```
In [38]: print("Coefficient:",svr.coef_)  
         print("Intercept:",svr.intercept_)
```

```
Coefficient: [[0.97475992 0.06513406 0.01905855]]  
Intercept: [30723.4595925]
```

```
In [39]: predicted_values_SVR = svr.predict(X_test)  
         predicted_values_SVR[0:10]
```

```
Out[39]: array([163836.57907645, 182175.35378806, 119725.7061669 ,  90266.47418753,  
                135272.5509011 ,  69766.64764945,  45211.86713187, 163849.04183737,  
                141635.16325919,  58268.79318778])
```

```
In [40]: print("Train score:",svr.score(X_train,y_train))  
         print("Test Score:",svr.score(X_test,y_test))
```

```
Train score: 0.9230196994607466  
Test Score: 0.8277782610994939
```

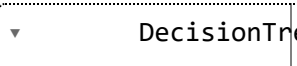
```
In [41]: MSE_SVR = mean_squared_error(y_test,predicted_values_SVR)  
         MAE_SVR = mean_squared_error(y_test,predicted_values_SVR)  
         RMSE_SVR = np.sqrt(MSE_ENR)
```

```
In [42]: SVR_score = r2_score(y_test,predicted_values_SVR)  
         SVR_score
```

```
Out[42]: 0.8277782610994939
```

## Decision Tree Regressor Algorithm

```
In [43]: from sklearn.tree import DecisionTreeRegressor  
DTR = DecisionTreeRegressor(random_state = 0)  
DTR.fit(X_train,y_train)
```

```
Out[43]:  DecisionTreeRegressor  
DecisionTreeRegressor(random_state=0)
```

```
In [44]: predicted_values_DTR = DTR.predict(X_test)  
predicted_values_DTR[0:10]
```

```
Out[44]: array([156991.12, 156991.12, 111313.02,  96479.51, 129917.04,  90708.19,  
                69758.98, 125370.37, 129917.04,  71498.49])
```

```
In [45]: print("Train score:",DTR.score(X_train,y_train))  
print("Test Score:",DTR.score(X_test,y_test))
```

```
Train score: 1.0  
Test Score: 0.8511432543536505
```


```
In [46]: MSE_DTR = mean_squared_error(y_test,predicted_values_DTR)  
MAE_DTR = mean_squared_error(y_test,predicted_values_DTR)  
RMSE_DTR = np.sqrt(MSE_DTR)
```

```
In [47]: DTR_score = r2_score(y_test,predicted_values_DTR)  
DTR_score
```

```
Out[47]: 0.8511432543536505
```

## Random Forrest Regressor Algorithm

```
In [48]: from sklearn.ensemble import RandomForestRegressor  
RFR = RandomForestRegressor(random_state = 0)  
RFR.fit(X_train,y_train)
```

```
Out[48]:   
RandomForestRegressor(random_state=0)
```

```
In [49]: predicted_values_RFR = RFR.predict(X_test)  
predicted_values_RFR[0:10]
```

```
Out[49]: array([146457.3604, 168421.306 , 115010.6291,  94594.7684, 132181.5484,  
                82692.8431,  60660.0496, 142125.2602, 136495.0184,  72659.7386])
```

```
In [50]: print("Train score:",RFR.score(X_train,y_train))  
print("Test Score:",RFR.score(X_test,y_test))
```

```
Train score: 0.9883903760207625  
Test Score: 0.9471039315439651
```

```
In [51]: MSE_RFR = mean_squared_error(y_test,predicted_values_RFR)  
MAE_RFR = mean_squared_error(y_test,predicted_values_RFR)  
RMSE_RFR = np.sqrt(MSE_RFR)
```

```
In [52]: RFR_score = r2_score(y_test,predicted_values_RFR)  
RFR_score
```

```
Out[52]: 0.9471039315439651
```

## Gradient Boosting Regressor Algorithm

```
In [53]: from sklearn.ensemble import GradientBoostingRegressor  
GBR = GradientBoostingRegressor(random_state = 0)  
GBR.fit(X_train,y_train)
```

```
Out[53]: GradientBoostingRegressor  
GradientBoostingRegressor(random_state=0)
```

```
In [54]: predicted_values_GBR = GBR.predict(X_test)  
predicted_values_GBR[0:10]
```

```
Out[54]: array([154481.2977341 , 168711.31499225, 112959.37189396,  97493.49024188,  
                132670.44103907,  87345.97091729,  63676.9309696 , 132957.99037655,  
                133497.08360501,  56543.51633607])
```

```
In [55]: print("Train score:",GBR.score(X_train,y_train))  
print("Test Score:",GBR.score(X_test,y_test))
```

```
Train score: 0.9999527462397211  
Test Score: 0.8806478515396146
```

```
In [56]: MSE_GBR = mean_squared_error(y_test,predicted_values_GBR)  
MAE_GBR = mean_squared_error(y_test,predicted_values_GBR)  
RMSE_GBR = np.sqrt(MSE_GBR)
```

```
In [57]: GBR_score = r2_score(y_test,predicted_values_GBR)  
GBR_score
```

```
Out[57]: 0.8806478515396146
```

We have applied 8 regression algorithms in the dataset .Those are

- **Linear Regression**
- **Ridge Regression**
- **Lasso Regression**



- **Elastic net Regression**
- **Support Vector Regressor**
- **Decision Tree Regressor**
- **Random Forest Regressor**
- **Gradient Boosting Regressor**

Lets make a tabular form of the following results that we have got so we would have a better explanation of which algorithm is to be considered.



```
In [58]: x = {"Linear Regression":lin.score(X_train,y_train),
            "Ridge Regression":ridge.score(X_train,y_train),
            "Lasso Regression":lasso.score(X_train,y_train),
            "Elastic net Regression":elastic_net.score(X_train,y_train),
            "Support Vector Regressor":svr.score(X_train,y_train),
            "Decision Tree Regressor":DTR.score(X_train,y_train),
            "Random Forest Regressor":RFR.score(X_train,y_train),
            "Gradient Boosting Regressor":GBR.score(X_train,y_train)}
y = {"Linear Regression":LR_score,
      "Ridge Regression":RR_score,
      "Lasso Regression":LAR_score,
      "Elastic net Regression":ENR_score,
      "Support Vector Regressor":SVR_score,
      "Decision Tree Regressor":DTR_score,
      "Random Forest Regressor":RFR_score,
      "Gradient Boosting Regressor":GBR_score}
z = {"Linear Regression": MSE_LR,
      "Ridge Regression": MSE_RR,
      "Lasso Regression": MSE_LAR,
      "Elastic net Regression":MSE_ENR,
      "Support Vector Regressor":MSE_SVR,
      "Decision Tree Regressor":MSE_DTR,
      "Random Forest Regressor":MSE_RFR,
      "Gradient Boosting Regressor":MSE_GBR}
u = {"Linear Regression":MAE_LR,
      "Ridge Regression":MAE_RR,
      "Lasso Regression":MAE_LAR,
      "Elastic net Regression":MAE_ENR,
      "Support Vector Regressor":MAE_SVR,
      "Decision Tree Regressor":MAE_DTR,
      "Random Forest Regressor":MAE_RFR,
      "Gradient Boosting Regressor":MAE_GBR}
v = {"Linear Regression":RMSE_LR,
      "Ridge Regression":RMSE_RR,
      "Lasso Regression":RMSE_LAR,
      "Elastic net Regression":RMSE_ENR,
      "Support Vector Regressor":RMSE_SVR,
      "Decision Tree Regressor":RMSE_DTR,
      "Random Forest Regressor":RMSE_RFR,
      "Gradient Boosting Regressor":RMSE_GBR}

k = pd.Series(x)
l = pd.Series(y)
```

```

m = pd.Series(z)
n = pd.Series(u)
v = pd.Series(v)

table = pd.DataFrame({"Train Score":k, "Accuracy":l, "MSE":m, "MAE":n, "RMSE":v})
table

```

Out[58]:

	Train Score	Accuracy	MSE	MAE	RMSE
<b>Linear Regression</b>	0.953967	0.926387	8.856874e+07	8.856874e+07	9411.096675
<b>Ridge Regression</b>	0.953967	0.926387	8.856874e+07	8.856874e+07	9411.096675
<b>Lasso Regression</b>	0.953967	0.926387	8.856874e+07	8.856874e+07	9411.096708
<b>Elastic net Regression</b>	0.953967	0.926387	8.856874e+07	8.856874e+07	9411.096677
<b>Support Vector Regressor</b>	0.923020	0.827778	2.072125e+08	2.072125e+08	9411.096677
<b>Decision Tree Regressor</b>	1.000000	0.851143	1.791003e+08	1.791003e+08	13382.837831
<b>Random Forest Regressor</b>	0.988390	0.947104	6.364310e+07	6.364310e+07	7977.662372
<b>Gradient Boosting Regressor</b>	0.999953	0.880648	1.436012e+08	1.436012e+08	11983.372954

- So from all these usecase of the algorithms we get to know that, **Linear Regression, Ridge regression, Lasso Regression, Elastic net Regression** are the algorithm gives almost the same prediction with the same accuracy score also efficiency of these algorithms are very much higher.
- The best efficiency we are getting from the Algorithm is **Random Forest Regressor**
- so we can select any of the Regression algorithm from the above mentioned algorithm for the prediction of the profit. The above table shows all the metrics that have been used in the algorithm.

## Visualization of the Actual vs Predicted Profit

```
In [59]: import matplotlib.pyplot as plt
y_test_flat = y_test.values.flatten()
fig, axs = plt.subplots(4, 2, figsize=(12, 15))

"""Linear Regression"""

predicted_values_flat = predicted_values.flatten()
axs[0,0].scatter(y_test_flat, predicted_values_flat, color='blue', marker='o', label='Predicted')
k_1, intercept = np.polyfit(y_test_flat, predicted_values_flat, 1)
x_line = np.linspace(min(y_test_flat), max(y_test_flat), 100)
y_line = k_1 * x_line + intercept
axs[0,0].plot(x_line, y_line, color='red', label='Linear Regression Line')
axs[0,0].set_xlabel('Actual Profit')
axs[0,0].set_ylabel('Predicted Profit')
axs[0,0].set_title('Actual vs Predicted Profit (Linear Regression)')
axs[0,0].legend()

"""Ridge Regression"""

predicted_values_flat_RR = predicted_values_RR.flatten()
axs[0,1].scatter(y_test_flat, predicted_values_flat_RR, color='grey', marker='o', label='Predicted')
k_1, intercept = np.polyfit(y_test_flat, predicted_values_flat_RR, 1)
x_line = np.linspace(min(y_test_flat), max(y_test_flat), 100)
y_line = k_1 * x_line + intercept
axs[0,1].plot(x_line, y_line, color='black', label='Ridge Regression Line')
axs[0,1].set_xlabel('Actual Profit')
axs[0,1].set_ylabel('Predicted Profit')
axs[0,1].set_title('Actual vs Predicted Profit (Ridge Regression)')
axs[0,1].legend()

"""Lasso Regression"""

predicted_values_flat_LAR = predicted_values_LAR.flatten()
axs[1,0].scatter(y_test_flat, predicted_values_flat_LAR, color='Violet', marker='o', label='Predicted')
k_1, intercept = np.polyfit(y_test_flat, predicted_values_flat_LAR, 1)
x_line = np.linspace(min(y_test_flat), max(y_test_flat), 100)
y_line = k_1 * x_line + intercept
axs[1,0].plot(x_line, y_line, color='midnightblue', label='Lasso Regression Line')
axs[1,0].set_xlabel('Actual Profit')
axs[1,0].set_ylabel('Predicted Profit')
axs[1,0].set_title('Actual vs Predicted Profit (Lasso Regression)')
axs[1,0].legend()
```

```
"""Elastic net Regression"""
```

```
predicted_values_flat_ENR = predicted_values_ENR.flatten()
axs[1,1].scatter(y_test_flat, predicted_values_flat_ENR, color='mediumseagreen', marker='o', label='Predicted')
k_1, intercept = np.polyfit(y_test_flat, predicted_values_flat_ENR, 1)
x_line = np.linspace(min(y_test_flat), max(y_test_flat), 100)
y_line = k_1 * x_line + intercept
axs[1,1].plot(x_line, y_line, color='darkslategrey', label='Elastic net Regression Line')
axs[1,1].set_xlabel('Actual Profit')
axs[1,1].set_ylabel('Predicted Profit')
axs[1,1].set_title('Actual vs Predicted Profit (Elastic net Regression)')
axs[1,1].legend()
```

```
"""Support Vector Regressor"""
```

```
predicted_values_flat_SVR = predicted_values_SVR.flatten()
axs[2,0].scatter(y_test_flat, predicted_values_flat_SVR, color='lightskyblue', marker='o', label='Predicted')
k_1, intercept = np.polyfit(y_test_flat, predicted_values_flat_SVR, 1)
x_line = np.linspace(min(y_test_flat), max(y_test_flat), 100)
y_line = k_1 * x_line + intercept
axs[2,0].plot(x_line, y_line, color='darkmagenta', label='Support Vector Regressor line')
axs[2,0].set_xlabel('Actual Profit')
axs[2,0].set_ylabel('Predicted Profit')
axs[2,0].set_title('Actual vs Predicted Profit (Support Vector Regressor)')
axs[2,0].legend()
```

```
"""Decision Tree Regressor"""
```

```
predicted_values_flat_DTR = predicted_values_DTR.flatten()
axs[2,1].scatter(y_test_flat, predicted_values_flat_DTR, color='cyan', marker='o', label='Predicted')
k_1, intercept = np.polyfit(y_test_flat, predicted_values_flat_DTR, 1)
x_line = np.linspace(min(y_test_flat), max(y_test_flat), 100)
y_line = k_1 * x_line + intercept
axs[2,1].plot(x_line, y_line, color='teal', label='Decision Tree Regressor Line')
axs[2,1].set_xlabel('Actual Profit')
axs[2,1].set_ylabel('Predicted Profit')
axs[2,1].set_title('Actual vs Predicted Profit (Decision Tree Regression)')
axs[2,1].legend()
```

```
"""Random Forest Regressor"""
```

```
predicted_values_flat_RFR = predicted_values_RFR.flatten()
axs[3,0].scatter(y_test_flat, predicted_values_flat_RFR, color='burlywood', marker='o', label='Predicted')
```

```
k_1, intercept = np.polyfit(y_test_flat, predicted_values_flat_RFR, 1)
x_line = np.linspace(min(y_test_flat), max(y_test_flat), 100)
y_line = k_1 * x_line + intercept
axs[3,0].plot(x_line, y_line, color='crimson', label='Random Forest Regressor Line')
axs[3,0].set_xlabel('Actual Profit')
axs[3,0].set_ylabel('Predicted Profit')
axs[3,0].set_title('Actual vs Predicted Profit (Random Forest Regressor)')
axs[3,0].legend()



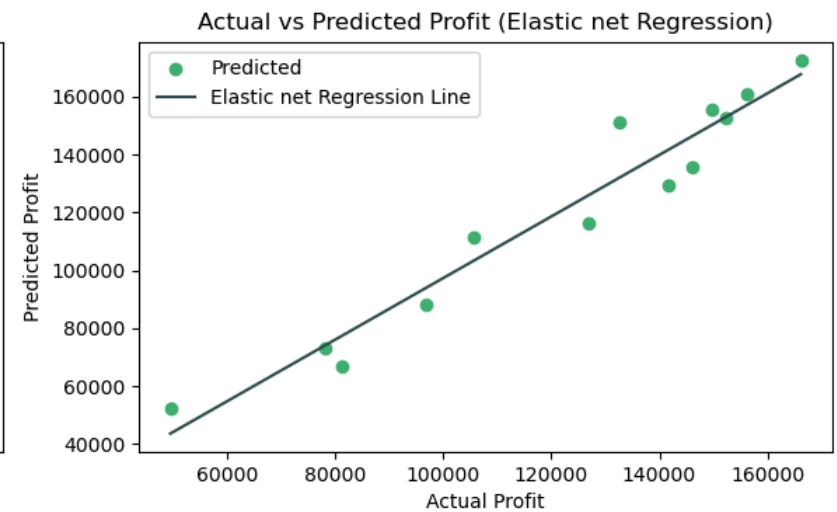



""""Gradient Boosting Regressor""""

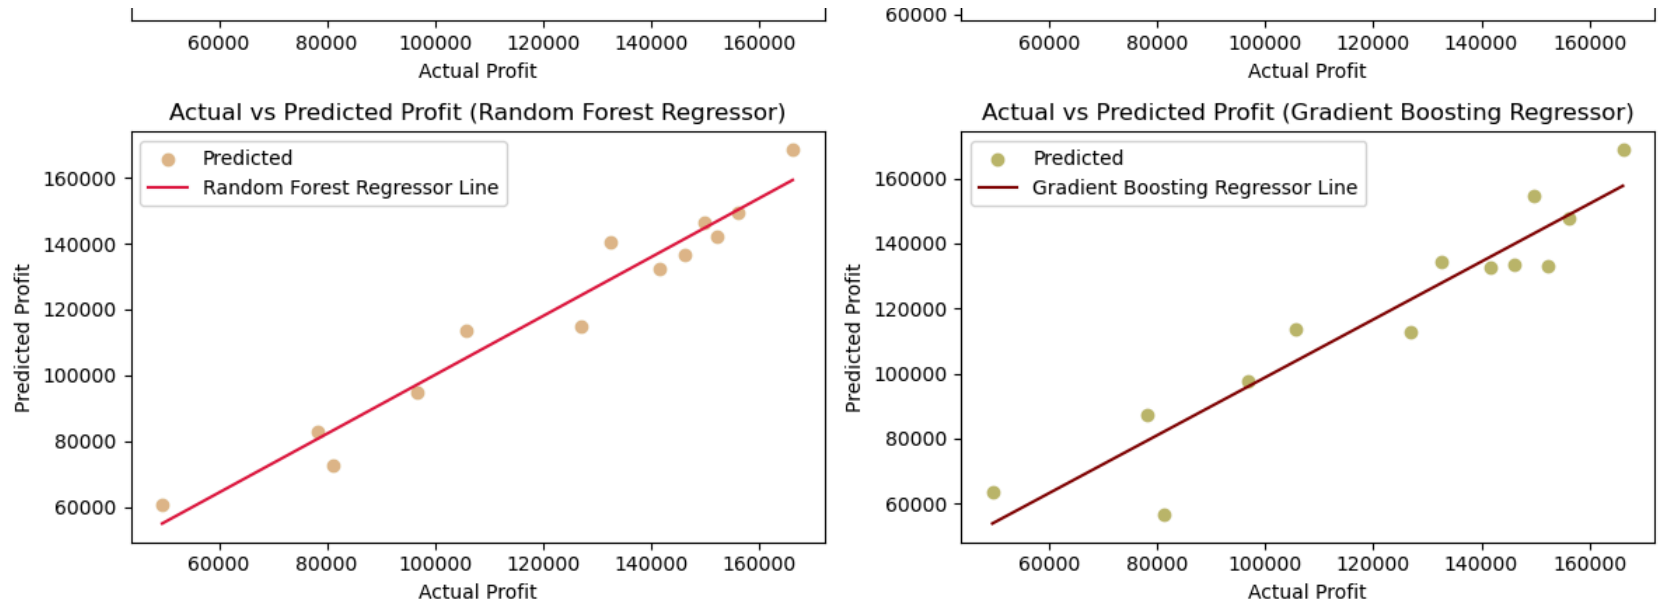
predicted_values_flat_GBR = predicted_values_GBR.flatten()
axs[3,1].scatter(y_test_flat, predicted_values_flat_GBR, color='darkkhaki', marker='o', label='Predicted')
k_1, intercept = np.polyfit(y_test_flat, predicted_values_flat_GBR, 1)
x_line = np.linspace(min(y_test_flat), max(y_test_flat), 100)
y_line = k_1 * x_line + intercept
axs[3,1].plot(x_line, y_line, color='maroon', label='Gradient Boosting Regressor Line')
axs[3,1].set_xlabel('Actual Profit')
axs[3,1].set_ylabel('Predicted Profit')
axs[3,1].set_title('Actual vs Predicted Profit (Gradient Boosting Regressor)')
axs[3,1].legend()

plt.tight_layout()
plt.show();
```









**Thats all for the Prediction of Profit for the dataset 50\_startups using different Regression Algorithm.**

In [ ]: